YAROSLAV TKACHENKO

Vancouver, BC | 604-230-8734 | sapiensy@gmail.com | sap1ens.com

Skills

Highly scalable and available service-oriented applications with **Scala**, **Java**, **Python** and **Node.js** stacks. Everything between API and infrastructure levels: starting from basic modelling and design, Domain-Driven back-end implementation (functional, OOP, actors), messaging, caching, background jobs, persistence with RDBMS and NoSQL, etc. to complex stateful cluster services, microservices, event sourcing, stream and batch processing. Big interest in **data pipelines** and **Big Data**.

Reliable reproducible **infrastructure-as-code** and **automation**. **Containers**. Twelve-Factor App adoration and **DevOps** excellence. Test automation junky.

Frequently **speaking** at tech conferences and conducting **workshops**.

Experience

Activision

Software Architect / *September 2018 – Present / Vancouver, Canada*

Senior Data Engineer / *March 2018 – September 2018 / Vancouver, Canada*

Senior Software Engineer at Demonware / *May 2017 – March 2018 / Vancouver, Canada*

Working on a Big Data platform for Activision games (primarily Call of Duty).

Initially started in the Data Pipeline team at Demonware, focusing on large-scale ingestion of data (Kafka). Worked on the pipeline architecture, infrastructure (Docker, Ansible, Terraform, AWS and on-premises), ingestion HTTP API (Python) and various consumers (Python, Java), including integration with Kafka Connect and stream processing with Kafka Streams.

Later transitioned to the Data Services team (Activision Central Tech). Working on a large-scale data pipeline (Kafka), ETL (Hadoop, Hive, Spark, Scala, Java), Data Lake / Warehouse (S3, Hive, Presto, Redshift) and various related services and APIs (Python, Java, MySQL, Cassandra, Redis). Running on AWS and leveraging Qubole.

In my new role as a Software Architect I'm responsible for architecting, designing and leading implementation efforts for the next-generation Big Data platform. Currently focused on developing realtime streaming ETL pipeline.

Mobify

Senior Software Engineer, Lead / *November 2016 - May 2017 / Vancouver, Canada*

Led the new Platform team, responsible for core Platform services like demo environments, content delivery, platform generator, location services and more. Heavily used Node.js, Python, Redis, Docker, AWS and Heroku.

Tried to invert the Conway's law and build unified Platform from a set of diverse products.

Designed and implemented first AMP (Accelerated Mobile Pages) prototype (ES6, Node.js, Redis).

Contributed to the realtime data pipeline (Scala, Kafka, Flink) and Progressive Mobile SDK (ES6, React).

Bench Accounting

Director of Engineering / *October 2015 - October 2016 / Vancouver, Canada*

Engineering Lead / *April 2014 - September 2015 / Vancouver, Canada*

Software Engineer / *September 2011 - March 2014 / Omsk, Russia (Remote) -> Vancouver, Canada*

Worked on a front-end and UI first year, rewriting multiple single page web apps, then focused more on API, back-end of the core services and integrations. Finally, I was responsible for the whole Platform, including infrastructure work, DevOps tools and practices, core back-end services, integrations, microservices architecture and more.

I always cared a lot about the projects we develop and products we ship. Everything from tiny UI elements to code quality and team motivation. I quickly became a leader for many initiatives and was promoted to the Engineering Lead. As an example, I led the transition from Java monolith app to Scala microservices.

Main projects and achievements included:

- Rewriting front-end apps from tangled jQuery to CoffeeScript, Backbone.js and RequireJS

- RESTful HTTP API for client and financial data (Scala/Jersey), as well as various APIs for other services and integrations (Scala/Spray)

- Core financial model design and implementation (Java/Spring, MySQL/Hibernate)

- Billing/subscriptions functionality via integration with Stripe (Java/Spring)

- Highly-available J2EE/Tomcat setup with zero-downtime deployment

- Scala/Akka microservice template, supporting tooling and libraries. Used for 10+ services

- Realtime eventing infrastructure that powers all notifications and chat messages (Scala/Akka, ActiveMQ/Camel, Postgres, PubNub)

- Various integrations including Salesforce, Zuora, FileThis, etc. (Scala/Akka)

- Parsers/scrapers for multiple financial institutions (Scala/Akka)

- End-to-end testing for the whole system (Selenium/Java and Intern.js/CoffeeScript later)

- Creating reproducible infrastructure environments for the microservices architecture (Ansible, AWS CloudFormation)

- Highly-available and auto-scalable set of microservices with zero-downtime deployment, unified monitoring and logging (Docker, AWS Elastic Beanstalk, Ansible, Jenkins, Datadog, Splunk)

In my role as a Director of Engineering I was responsible for managing multiple teams, working on Platform, infrastructure, API and integrations, as well as participating in technical discussions, contributing to architecture and mentoring (but still hands-on).

Purebulk

Senior Web Developer / *March 2010 - July 2011 / Omsk, Russia (Remote)*

Big e-commerce website (up to 20k+ active users daily), responsibilities included:

- Maintenance and feature development (PHP, MySQL)

- Heavy front-end development using various JavaScript frameworks

- Zencart -> Magento migration

Freelance & Digital Agencies

Web Developer / *January 2007 - February 2010 / Omsk, Russia*

Full-stack (HTML, CSS, JavaScript, PHP, MySQL) web development for various clients.

Education

Omsk State Technical University

2007 - 2012 / Omsk, Russia

Master's Degree (equivalent), Software Engineering and System Analysis.

Training & Courses

Serverless Workshop / *September 2017 / Amazon Web Services*

Implementing Domain-Driven Design workshop / *March 2016 / by Vaughn Vernon*

Deploying and scaling applications with Docker / *September 2016 / Strange Loop*

Writing Microservices with Lagom / *September 2016 / Strange Loop*

Introduction to Rust / *September 2016 / Strange Loop*

Introduction to Machine Learning Using Spark / *May 2015 / Polyglot Unconference*

Real World Programming With Haskell / *May 2014 / Polyglot Unconference*

Certifications

Big Data Analysis with Scala and Spark / *August 2017 / Coursera*

Principles of Reactive Programming / *May 2015 / Coursera*

Functional Programming Principles in Scala / *June 2014 / Coursera*